



Deep Learning Overview

Images and Video



Agenda

- I. Deep learning in engineering and science
- II. Developing a deep learning solution in MATLAB
- III. MathWorks deep learning support




Agenda

- I. Deep learning in engineering and science
- II. Developing a deep learning solution in MATLAB
- III. MathWorks deep learning support

Deep learning is a key technology driving the AI megatrend

ARTIFICIAL INTELLIGENCE

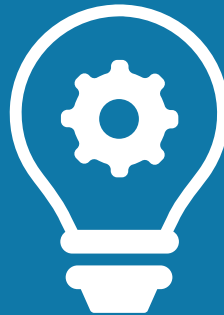
Any technique that enables machines to mimic human intelligence



1950s

MACHINE LEARNING

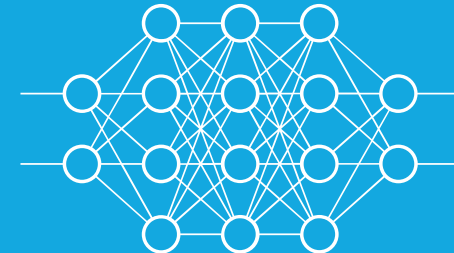
Statistical methods that enable machines to “learn” tasks from data without explicitly programming



1980s

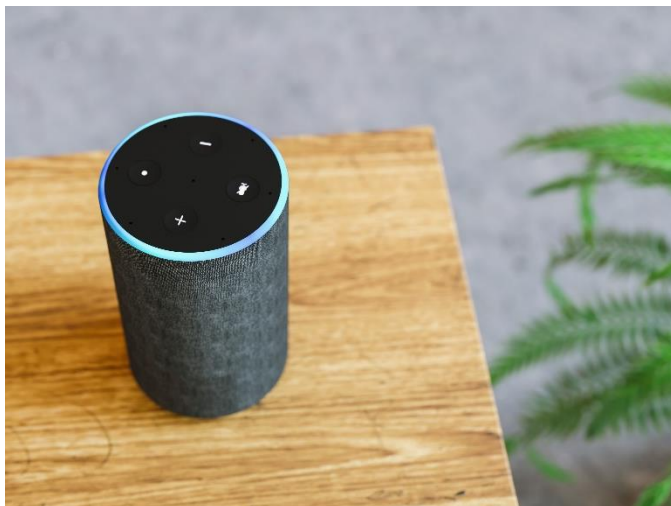
DEEP LEARNING

Neural networks with many layers that learn representations and tasks “directly” from data



2010s

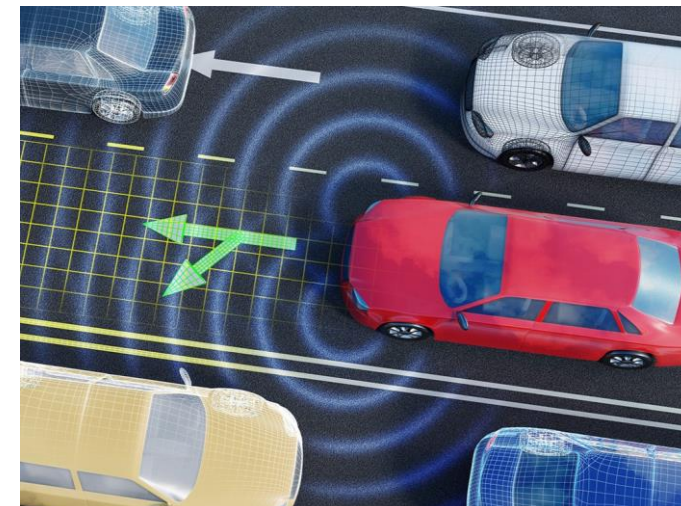
Deep learning is part of our everyday lives



Speech Recognition



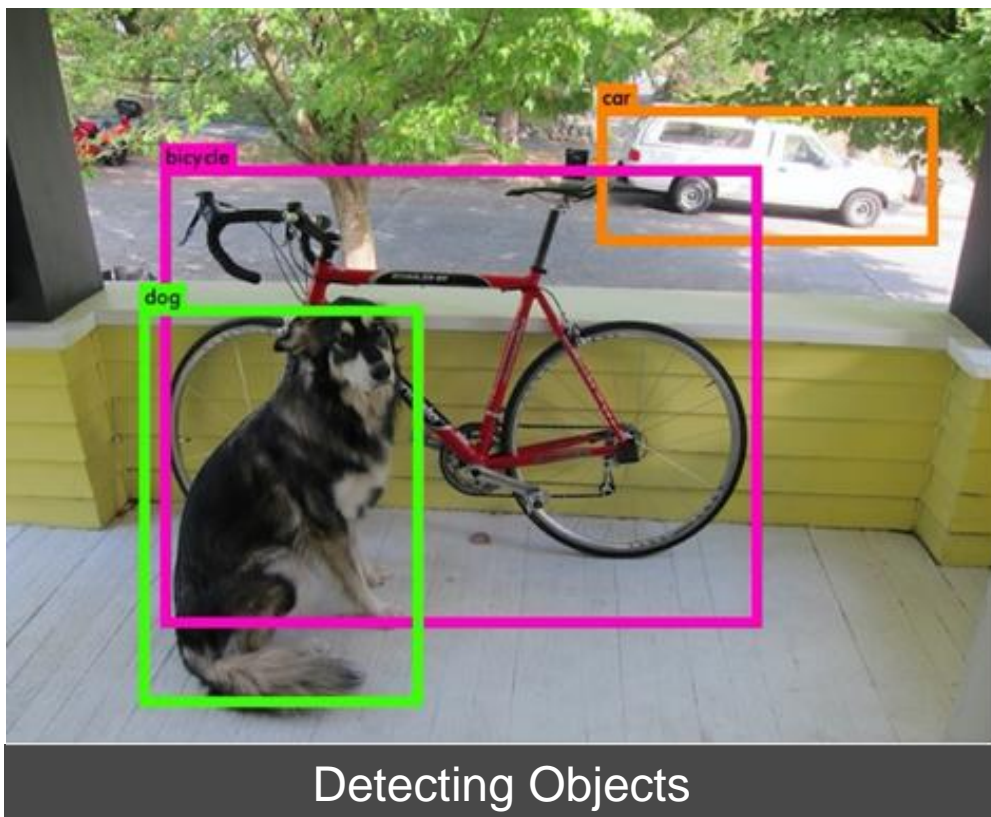
Face Detection



Automated Driving

Deep learning applications: mainstream vs. engineering

Mainstream



Engineering and Science

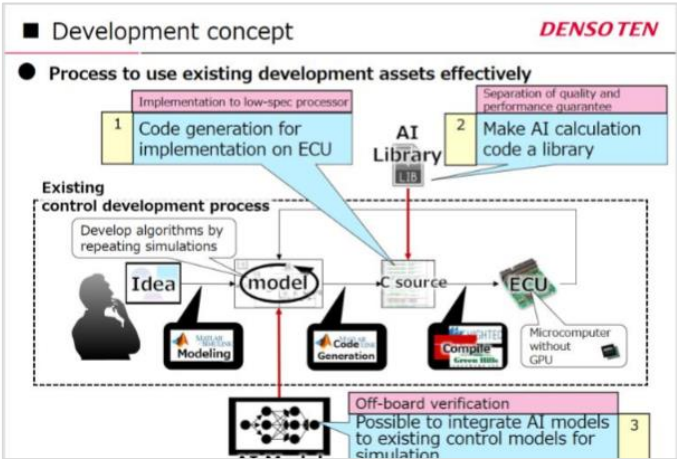


Deep Learning Detection


MATLAB Deep Learning used in Industry



Automatic Defect Detection
Airbus




ECU Vehicle Control
Denso

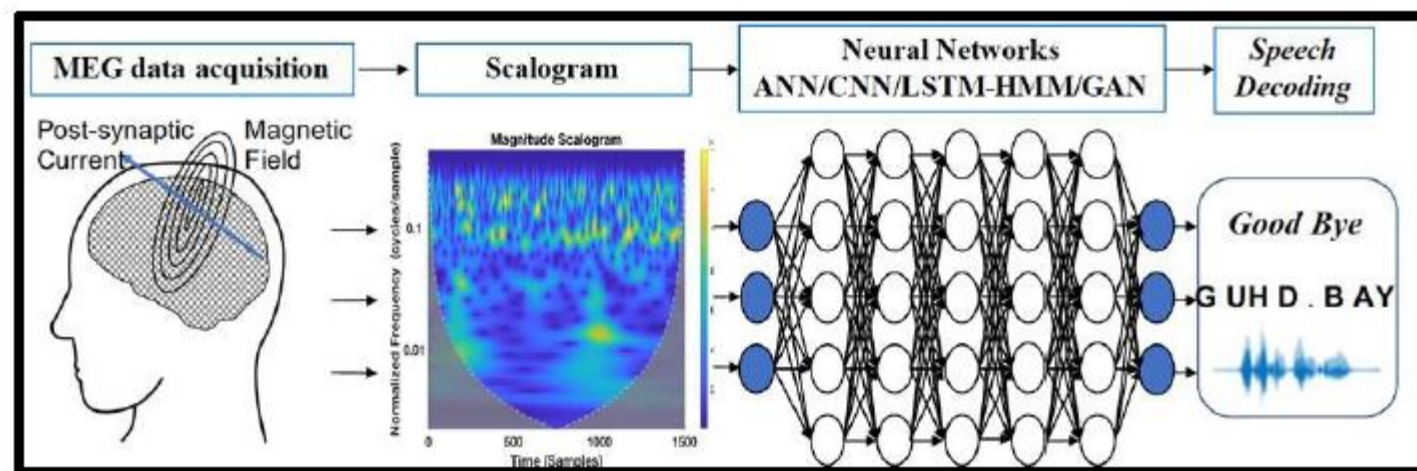


Seismic Event Detection
Shell

MATLAB Deep Learning used in Research



Predicting gastrointestinal cancer (July 2019)



Converting brain waves to speech to help ALS patients communicate (Nov 2019)

Evolution of Deep Learning in MATLAB

2016

CNN's
Pretrained Models
Caffe Importer

2017

Name Change

- Neural Network Toolbox to Deep Learning Toolbox

Algorithms

- LSTM's
- Directed Acyclic Graphs
- Multi-GPU Training

Code Generation

- GPU Coder

Apps

- Image Labeler

Interoperability

- TensorFlow-Keras Importer

2018

Examples

- Signal Processing
- Audio
- Text Analytics

Algorithms

- Wavelet Scattering

Code Generation

- MATLAB Coder C++

Apps

- Deep Network Designer
- Video Labeler
- Audio Labeler

Interoperability

- ONNX Support

2019

- Reinforcement Learning

Algorithms

- Automatic Differentiation
- Custom Training Loops
- Weight Sharing
- Big Image

Examples

- GANs
- Siamese Network
- Autoencoders
- 3-D support
- **Explainable AI**
 - Occlusion
 - Grad-CAM

Code Generation

- MATLAB Coder (ARM)

Apps

- Signal Labeler

2020

- Deep Learning Data Sets

Apps

- Experiment Manager
- Lidar Labeler

Examples

- 5G / Wireless
- Lidar

- **Over 200+ examples!**

Algorithms

- Point Cloud

Explainable AI

- **Lime**

Code Generation

- Quantization
- Deep Learning HDL Coder

Model Based Design

- Deep Learning in Simulink

Applications of deep learning for images and video

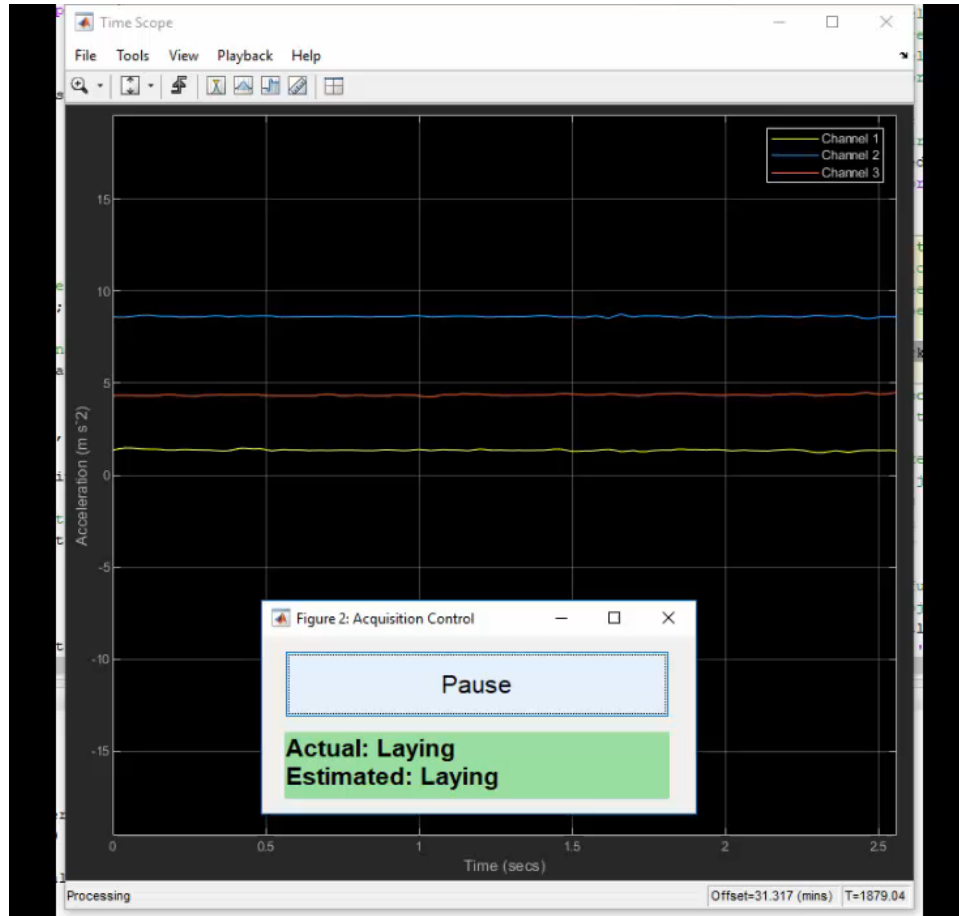


YOLO v2 (You Only Look Once)

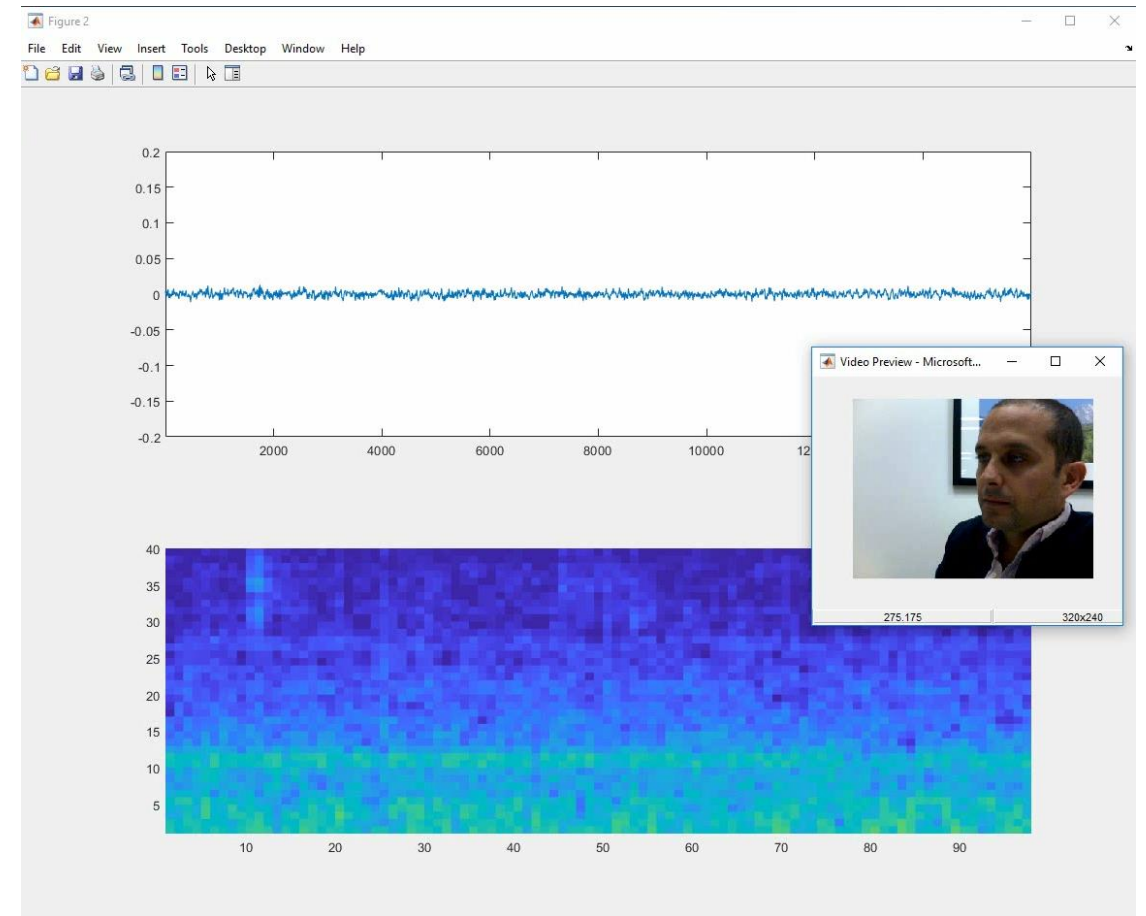


Semantic Segmentation using SegNet

Applications of deep learning for signal processing




Signal Classification using LSTMs




Speech Recognition using CNNs

Applications of reinforcement learning




Teach a car to navigate traffic




Train a Quadruped robot to walk

Why MATLAB & MathWorks for Deep Learning?

Domain-specialized workflows
for **engineering and science**



Multi-platform **deployment** of
full applications and systems




People



Platform productivity




Interoperability with
TensorFlow and PyTorch





is a **Leader** in the Gartner Magic Quadrant for 2020 Data Science and Machine Learning Platforms

Figure 1. Magic Quadrant for Data Science and Machine Learning Platforms



Source: Gartner (February 2020)

*Gartner Magic Quadrant for Data Science and Machine Learning Platforms, Peter Krensky, Erick Brethenoux, Jim Hare, Carlie Idoine, Alexander Linden, Svetlana Sicular, 11 February 2020 .

This graphic was published by Gartner, Inc. as part of a larger research document and should be evaluated in the context of the entire document. The Gartner document is available upon request from MathWorks. Gartner does not endorse any vendor, product or service depicted in its research publications, and does not advise technology users to select only those vendors with the highest ratings or other designation. Gartner research publications consist of the opinions of Gartner's research organization and should not be construed as statements of fact. Gartner disclaims all warranties, express or implied, with respect to this research, including any warranties of merchantability or fitness for a particular purpose.



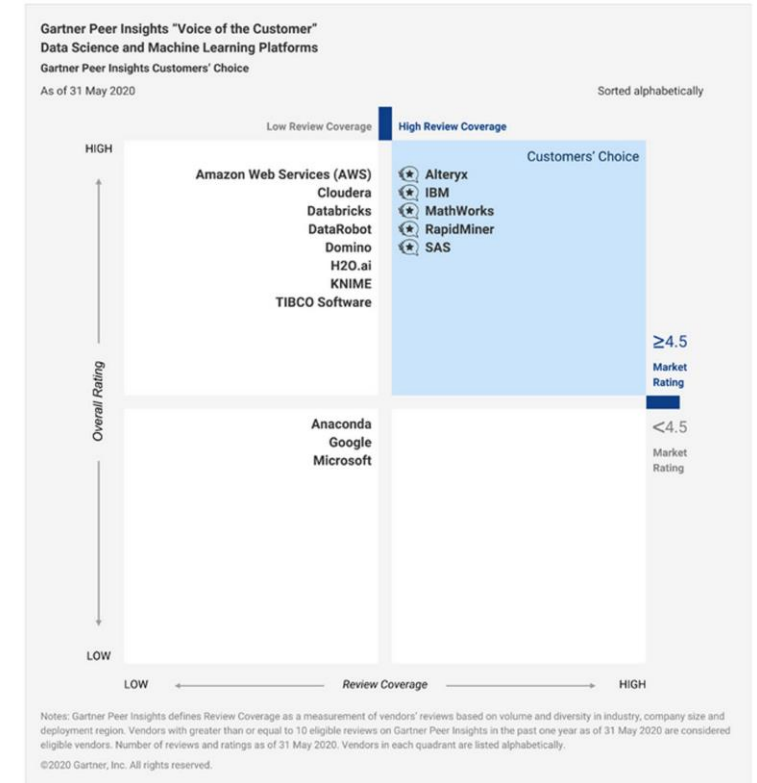
A 2020 Gartner Peer Insights Customers' Choice for Data Science and Machine Learning Platforms



is named a 2020 Gartner Peer Insights Customers' Choice for Data Science and Machine Learning Platforms.



Figure 1. Gartner Peer Insights "Voice of the Customer" Data Science and Machine Learning Platforms Customers' Choice



Source: Gartner (July 2020)

Disclaimer: Gartner, Gartner Peer Insights 'Voice of the Customer': Data Science and Machine Learning Platforms, July 2020. This graphic was published by Gartner, Inc. as part of a larger research document and should be evaluated in the context of the entire document. The Gartner document is available upon request from MathWorks.

The Gartner Peer Insights Customers' Choice badge is a trademark and service mark of Gartner, Inc., and/or its affiliates, and is used herein with permission. All rights reserved. Gartner Peer Insights Customers' Choice constitute the subjective opinions of individual end-user reviews, ratings, and data applied against a documented methodology; they neither represent the views of, nor constitute an endorsement by, Gartner or its affiliates.



AI-driven system design

Data Preparation



Data cleansing and preparation



Human insight



Simulation-generated data

AI Modeling



Model design and tuning



Hardware accelerated training



Interoperability

Simulation & Test



Integration with complex systems



System simulation



System verification and validation

Deployment



Embedded devices



Enterprise systems



Edge, cloud, desktop



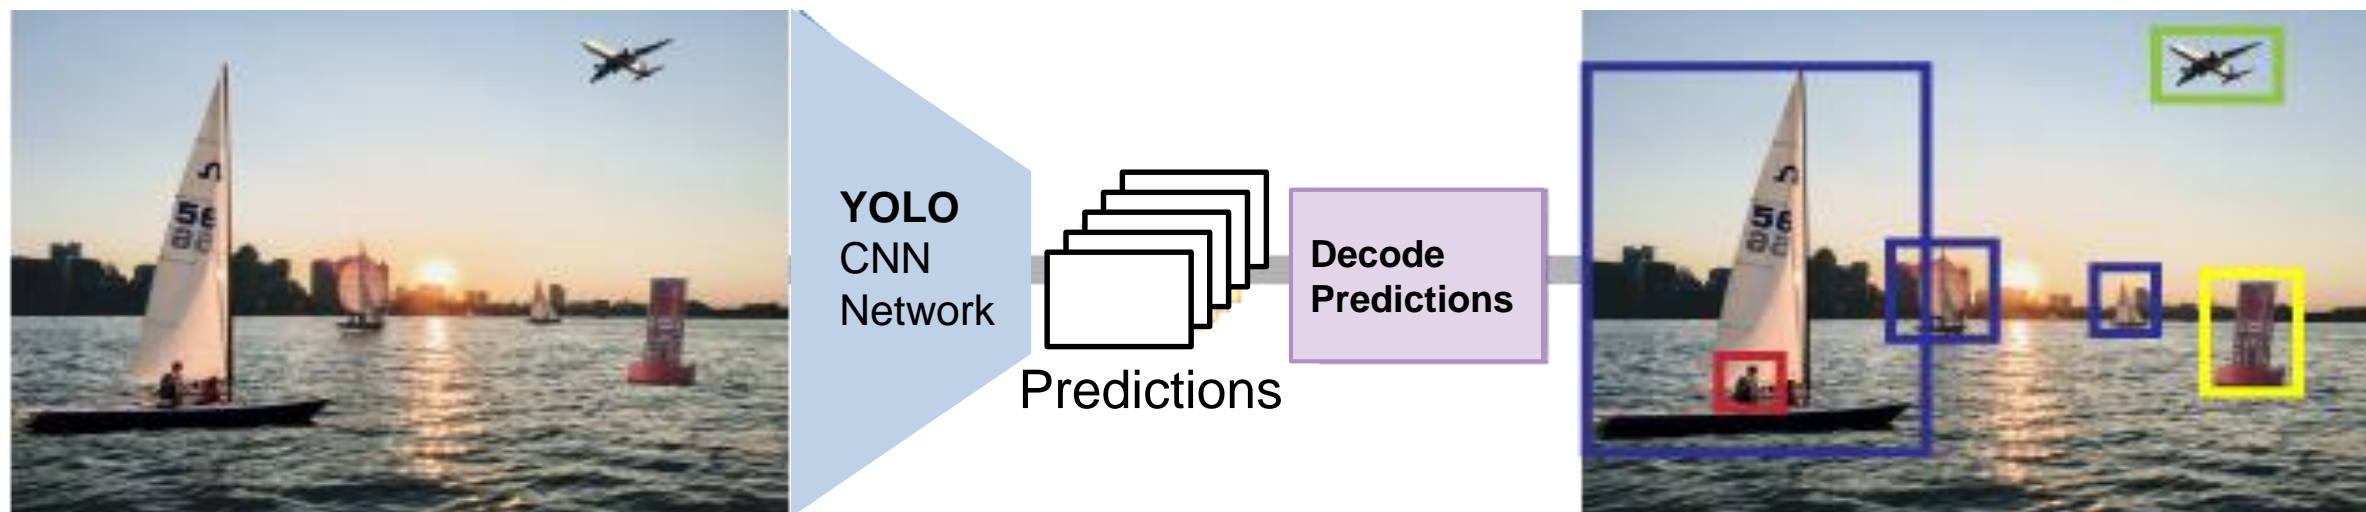
Agenda

- I. Deep learning in engineering and science
- II. Developing a deep learning solution in MATLAB
- III. MathWorks deep learning support

Featured Example: Detecting Objects with YOLO v2

Build, test, and deploy a deep learning solution that can detect objects in images and video.

- [You Only Look Once](#)
- Real-time object detector
- Autonomous driving, traffic monitoring
- 1000x faster than R-CNN



Data preparation represents most of your AI effort...


Transforming raw data for useful modeling and analysis is a critical step.

Data Preparation

 Data cleansing and preparation

 Human insight

 Simulation-generated data



Source: Andrej Karpathy slide from TrainAI 2018

Spend less time preprocessing and labeling data

Synchronize disparate time series, filter noisy signals, automate labeling of video, and more.

Data Preparation




Data cleansing and preparation



Human insight




Simulation-generated data



Use labeling apps for deep learning workflows like semantic segmentation

Data Preparation Demo



Open Script Part 1

Data Preparation




Data cleansing and
preparation



Human insight





Simulation-
generated data




Start with a complete set of algorithms and pre-built models

AI Modeling

 Model design and tuning

 Hardware accelerated training

 Interoperability

Algorithms

Machine learning

Trees, Naïve Bayes, SVM...

Deep learning

CNNs, GANs, LSTM, MIMO...

Reinforcement learning

DQN, A2C, DDPG...

Regression

Linear, nonlinear, trees...

Unsupervised learning

K-means, PCA, GMM...

Predictive maintenance

RUL models, condition indicators...

Bayesian optimization

Pre-built models

Image classification models

AlexNet, GoogLeNet, VGG,
SqueezeNet, ShuffleNet, ResNet,
DenseNet, Inception...

Reference examples

Object detection

Vehicles, pedestrians, faces...

Semantic segmentation

Roadway detection, land cover
classification, tumor detection...

Signal and speech processing

Denoising, music genre recognition,
keyword spotting, radar waveform
classification...

...and more...

Increase productivity using Apps for design and analysis

Use MATLAB Apps to design deep learning networks, explore a wide range of classifiers, train regression models, train an optical character recognition model, and more.

AI Modeling




Model design and tuning



Hardware accelerated training



Interoperability





Deep Network Designer app to build, visualize, and edit deep learning networks

Trial	Status	Progress	Elapsed Time	myInitialLearn...	convFilterSize	Training Accu...	Training Loss	Validation Ac...
1	Complete	100.0%	0 hr 0 min 16 sec	1.0000e-6	3.0000	12.5000	2.4441	10.
2	Complete	100.0%	0 hr 0 min 15 sec	1.0000e-5	3.0000	25.7813	2.1228	20.
3	Complete	100.0%	0 hr 0 min 14 sec	0.0001	3.0000	64.8438	1.0878	42.
4	Complete	100.0%	0 hr 0 min 16 sec	0.0005	3.0000	90.6250	0.4648	49.
5	Complete	100.0%	0 hr 0 min 15 sec	1.0000e-6	4.0000	11.7188	2.4967	6.
6	Complete	100.0%	0 hr 0 min 15 sec	1.0000e-5	4.0000	23.4375	2.1213	14.
7	Complete	100.0%	0 hr 0 min 17 sec	0.0001	4.0000	72.6563	1.0283	39.
8	Running	30.7%	0 hr 0 min 4 sec	0.0005	4.0000			
9	Queued	0.0%		1.0000e-6	5.0000			
10	Queued	0.0%		1.0000e-5	5.0000			
11	Queued	0.0%		0.0001	5.0000			
12	Queued	0.0%		0.0005	5.0000			
13	Queued	0.0%		1.0000e-6	6.0000			
14	Queued	0.0%		1.0000e-5	6.0000			
15	Queued	0.0%		0.0001	6.0000			
16	Queued	0.0%		0.0005	6.0000			


Experiment Manager app to manage multiple deep learning experiments, analyze and compare results and code


Hardware acceleration and scaling are critical for training


MATLAB accelerates AI training on GPUs, cloud, and datacenter resources without specialized programming.



AI Modeling

 Model design and tuning


 Hardware accelerated training


 Interoperability


Performance is Continuously Improving

Development continuously measures, tracks, and improves performance for training and inference

AI Modeling

 Model design and tuning


 Hardware accelerated training

 Interoperability

Training

R2020a-R2020b

Multi-GPU 1.6x speedup



Windows 10, Intel® Xeon® E5-2623 v4 @ 2.60 GHz
NVIDIA® Titan V 12GB GPUs

Inference

R2018b-R2020b

GPU 2.8x speedup


CPU 3x speedup


ResNet-50, Batch Size 32, Intel® Xeon® CPU 3.6 GHz – Titan V. NVIDIA


MATLAB interoperates with other frameworks

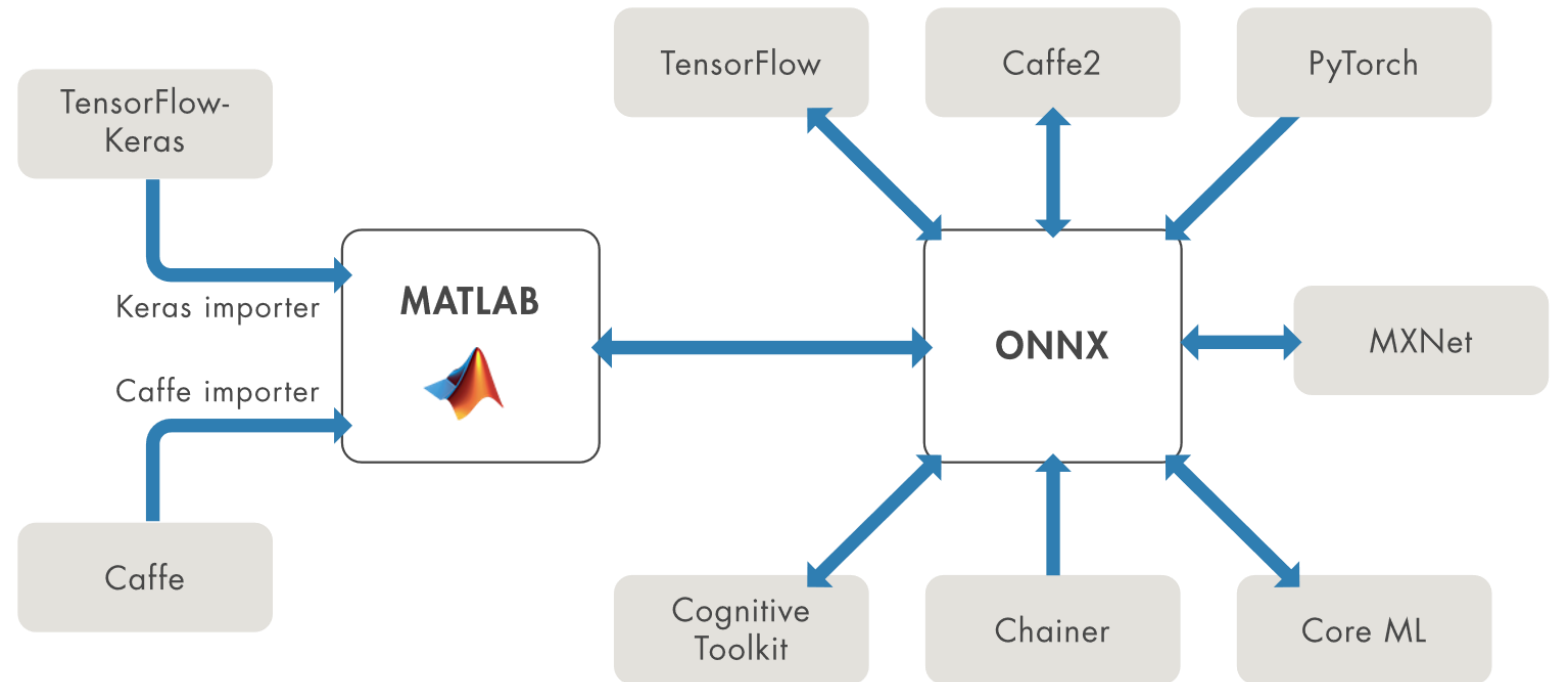
Supports ONNX and can exchange models with PyTorch, TensorFlow, and other frameworks.

AI Modeling


 Model design and tuning

 Hardware accelerated training

 Interoperability



Modeling Demo



**Open Script
Part 2**

AI Modeling




Model design and
tuning



Hardware
accelerated training



Interoperability



Models need to exist within a complete system

In automated driving systems, AI for perception must integrate with algorithms for path planning, braking, acceleration, and other controls.

Simulation & Test

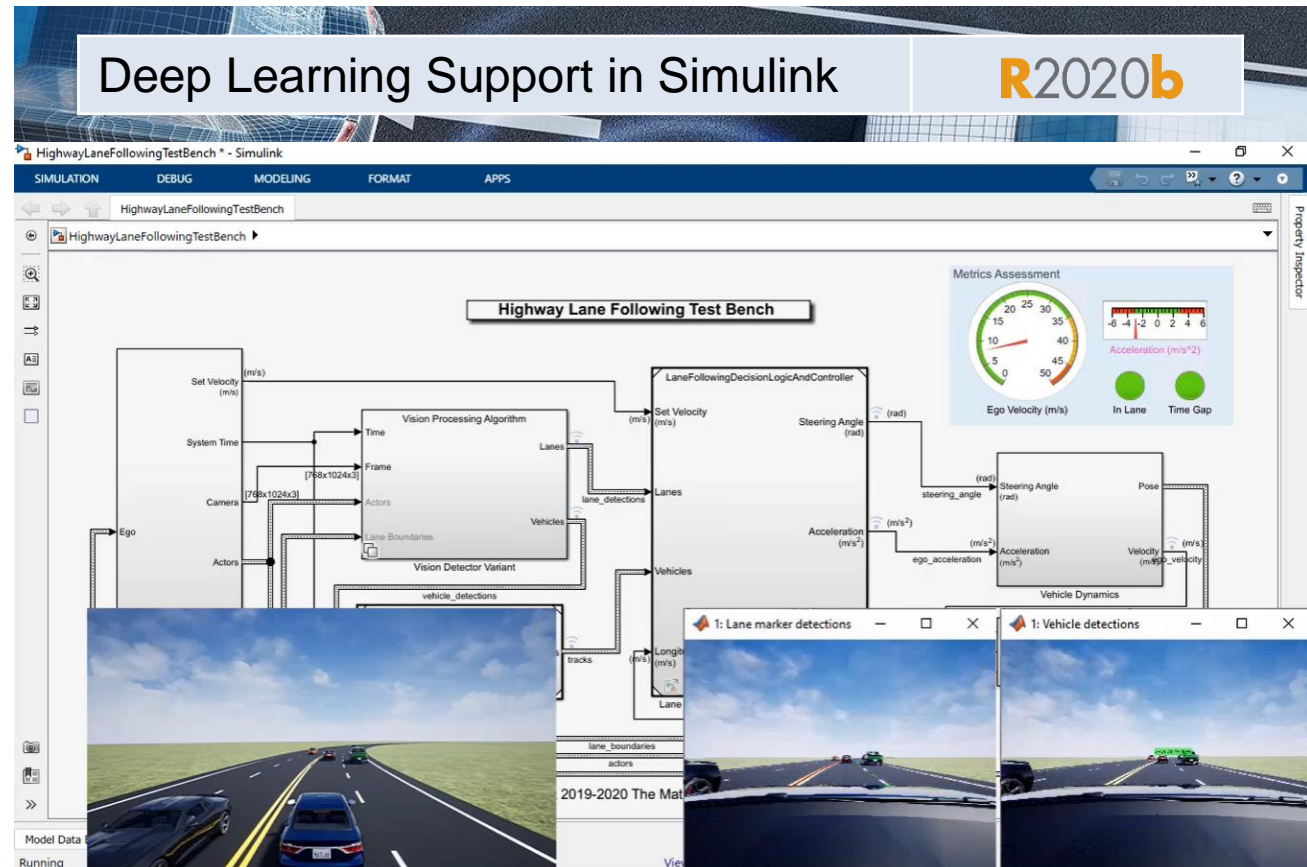


Integration with complex systems



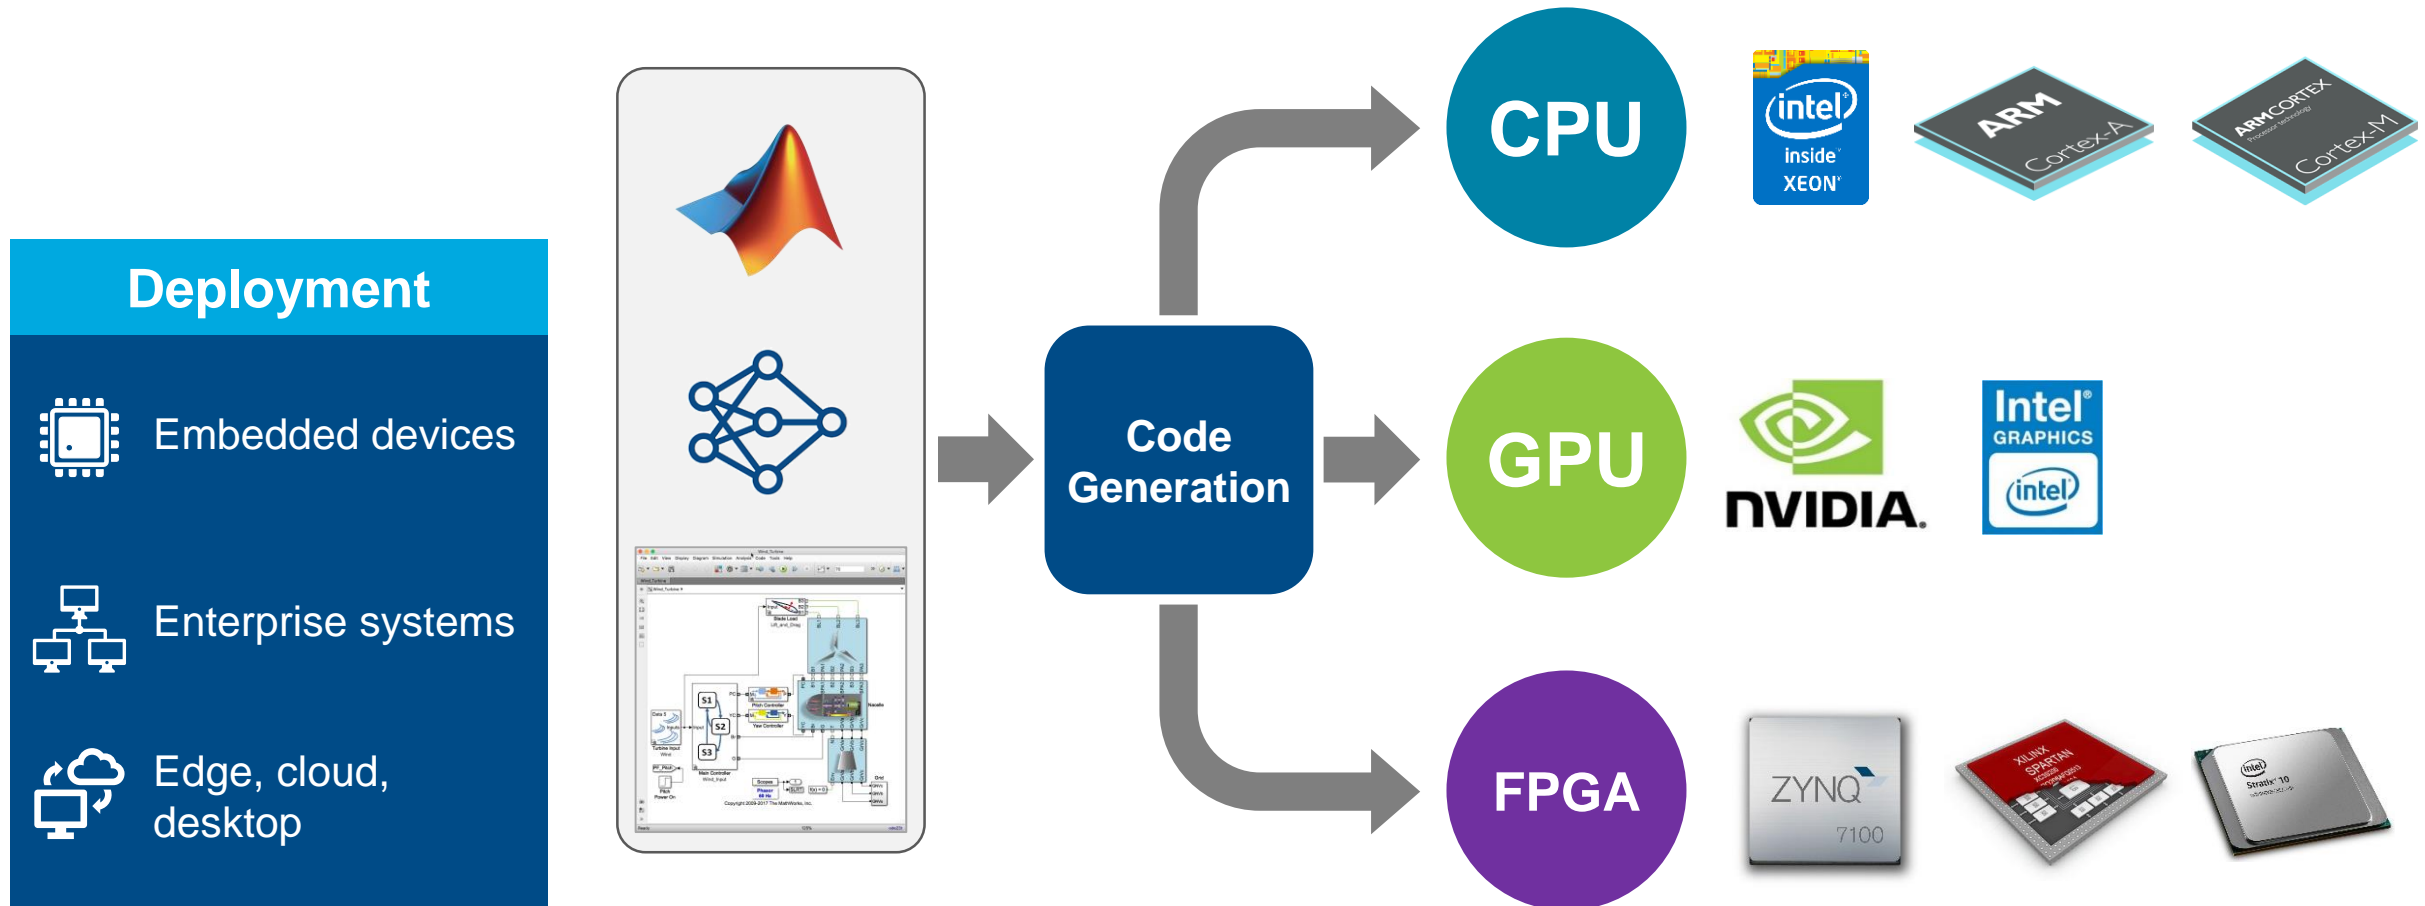
System simulation

— x System verification
— ✓ and validation

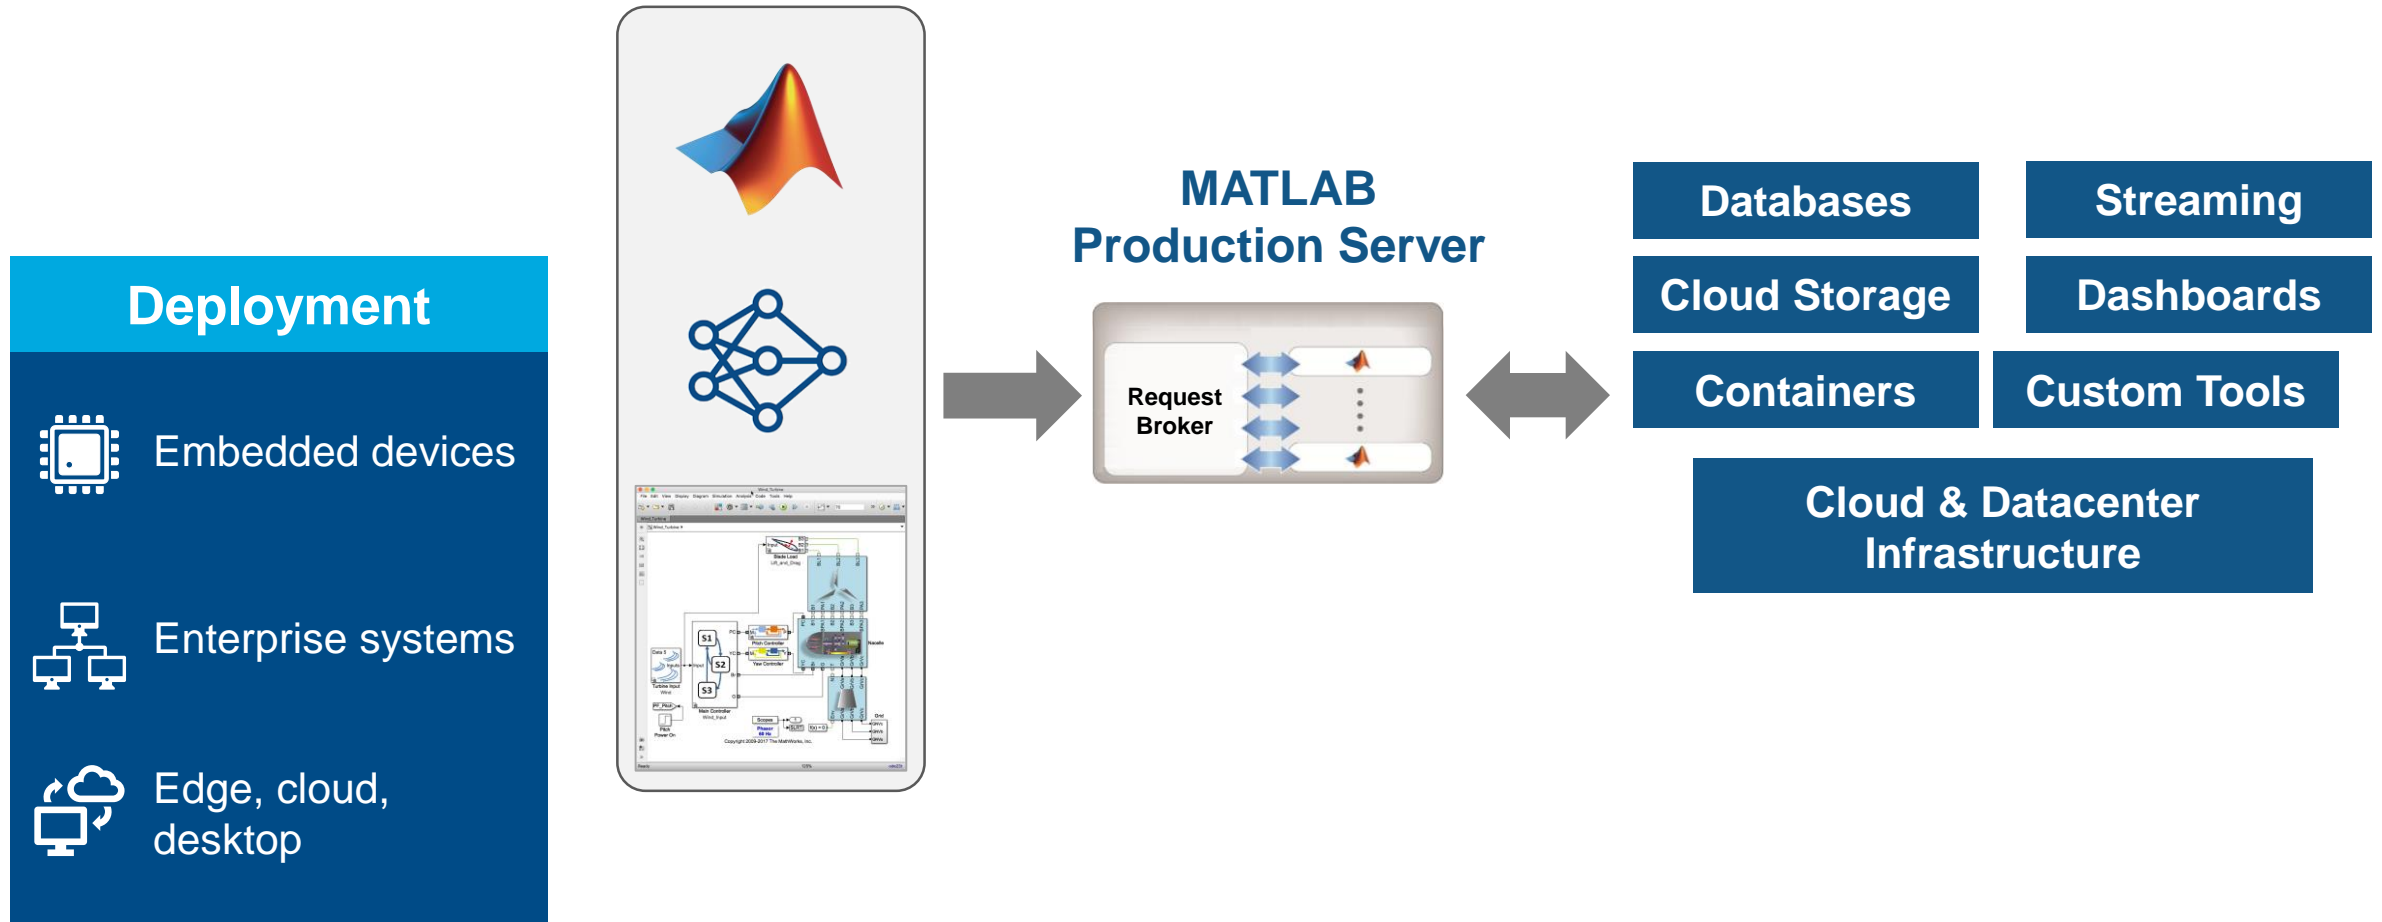


Deploy to any processor with best-in-class performance


AI models in MATLAB and Simulink can be deployed on embedded devices, edge devices, enterprise systems, the cloud, or the desktop.



Deploy to enterprise IT infrastructure



Deployment Demo



Open Script
Part 3

Deployment



Embedded devices



Enterprise systems



Edge, cloud,
desktop

The screenshot displays the MATLAB Cody environment. The left sidebar shows the 'MATLAB SOURCE' pane with a file list containing 'yolov2_detect.m', 'fx yolov2_detect' (marked GPU), 'cudnnApi.p', and 'fx cudnnApi_register'. Below it is the 'GENERATED CODE' pane with tabs for 'Files' and 'GPU Kernels', showing a tree of source files for the generated MEX function.

The main editor shows the MATLAB function 'yolov2_detect' with the following code:

```
1 function outImg = yolov2_detect(in)
2
3 % Copyright 2018-2019 The MathWorks, Inc.
4
5 persistent yolov2obj;
6
7 if isempty(yolov2obj)
8     yolov2obj = coder.loadDeepLearningNetwork('Yolov2UsingResNet50_ONNX.mat');
9 end
10
11 % pass in input
12 [bboxes,~,labels] = yolov2obj.detect(in,'Threshold',0.5);
13
14 % convert categorical labels to cell array of character vectors for MATLAB
15 % execution
16 if coder.target('MATLAB')
17     labels = cellstr(labels);
18 end
19
20 % Annotate detections in the image.
21 outImg = insertObjectAnnotation(in,'rectangle',bboxes,labels);
22
```

The bottom status bar shows a 'SUMMARY' tab with the message 'Code generation successful'. It includes the following details:

- Generated on: 17-Sep-2019 14:21:46
- Build type: MEX Function
- Output file: C:\Users\shmitra\Work\Deep_Learning\Seminar\19b\ResNetImportYolov2\HelperFilesAndFunctions\yolov2_detect_mex.mexw64
- Processor: Generic->MATLAB Host Computer



Agenda

- I. Deep learning in engineering and science
- II. Developing a deep learning solution in MATLAB
- III. MathWorks deep learning support



MathWorks is your Deep Learning partner



The Platform

MATLAB, Simulink, and over 100 add-on products for specialized applications



Your People

Helping you build an agile workforce today and preparing tomorrow's engineers



Our Expertise

From onboarding and implementation to solving advanced engineering challenges

MathWorks Engineering Support



Training



Guided Evaluations



Onsite Workshops




Consulting




Technical Support

Why MATLAB & MathWorks for Deep Learning?

Domain-specialized workflows
for **engineering and science**



Multi-platform **deployment** of
full applications and systems




People



Platform productivity





Interoperability with
TensorFlow and PyTorch



Further Learning & Teaching

- [Deep Learning Onramp](#)
 - 2 hr online tutorial
- Deep Learning Workshop
 - 3 hr hands on session
 - Contact us to schedule
- [Deep Learning Training](#)
 - 16 hr in depth course
 - Online or Instructor Lead
- [Teaching Deep Learning with MATLAB](#)
 - Curriculum support



Where to find this content on GitHub

- <https://tinyurl.com/deeplearningmatlabimage>



Learn Git and GitHub without any code!
Using the Hello World guide, you'll start a branch, write comments, and open a pull request.

[Read the guide](#)

matlab-deep-learning / Object-Detection-Using-YOLO-v2-Deep-Learning

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

MATLAB example of deep learning based object detection using Yolo v2.

matlab deep-learning example yolo-v2 object-detection Manage topics

9 commits 1 branch 0 packages 0 releases 1 contributor View license

Branch: master New pull request Create new file Upload files Find file Clone or download

		Latest commit 147b223 2 hours ago
CodeGeneration	Initial push to GitHub	19 hours ago
HelperFiles	Initial push to GitHub	19 hours ago
Images	Updated README.md	18 hours ago
SavedModels	Initial push to GitHub	19 hours ago
resources/project	Initial push to GitHub	19 hours ago
Deep Learning Overview For Images.pdf	Initial push to GitHub	19 hours ago
LICENSE	Initial push to GitHub	19 hours ago
Part01_DataPreparation.mlx	Initial push to GitHub	19 hours ago
Part02_Modeling.mlx	Updated README.md	18 hours ago
Part01_Deployment.mlx	Initial push to GitHub	19 hours ago
README.md	updated README.md	2 hours ago
YOLOv2ObjectDetection.prj	Initial push to GitHub	19 hours ago

README.md

Object Detection Example using Yolo v2 Deep Learning

This demo shows the full deep learning workflow for an example using image data in MATLAB. In it we use deep learning based object detection using Yolo v2 to identify vehicles of interest in a scene.

A photograph of a dark-colored SUV parked on a street. A yellow bounding box is drawn around the car, and the text "0.8185" is displayed below it, indicating the model's confidence in the detection.